

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from sklearn import preprocessing, svm
6 from sklearn.model_selection import train_test_split
7 from sklearn.linear_model import LinearRegression
```

In [2]:

```
1 dt=pd.read_csv(r"C:\Users\HP\Downloads\data.csv")
2 dt
```

Out[2]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	\
0	2014-05-02 00:00:00	3.130000e+05	3.0	1.50	1340	7912	1.5	0	
1	2014-05-02 00:00:00	2.384000e+06	5.0	2.50	3650	9050	2.0	0	
2	2014-05-02 00:00:00	3.420000e+05	3.0	2.00	1930	11947	1.0	0	
3	2014-05-02 00:00:00	4.200000e+05	3.0	2.25	2000	8030	1.0	0	
4	2014-05-02 00:00:00	5.500000e+05	4.0	2.50	1940	10500	1.0	0	
...	...	...	...	...	...	...	...	...	
4595	2014-07-09 00:00:00	3.081667e+05	3.0	1.75	1510	6360	1.0	0	
4596	2014-07-09 00:00:00	5.343333e+05	3.0	2.50	1460	7573	2.0	0	
4597	2014-07-09 00:00:00	4.169042e+05	3.0	2.50	3010	7014	2.0	0	
4598	2014-07-10 00:00:00	2.034000e+05	4.0	2.00	2090	6630	1.0	0	
4599	2014-07-10 00:00:00	2.206000e+05	3.0	2.50	1490	8102	2.0	0	

4600 rows × 18 columns



In [3]:

```
1 dt=dt[['sqft_living','sqft_lot']]
2 dt.columns=['Liv','Lot']
```

In [4]:

```
1 dt.head(10)
```

Out[4]:

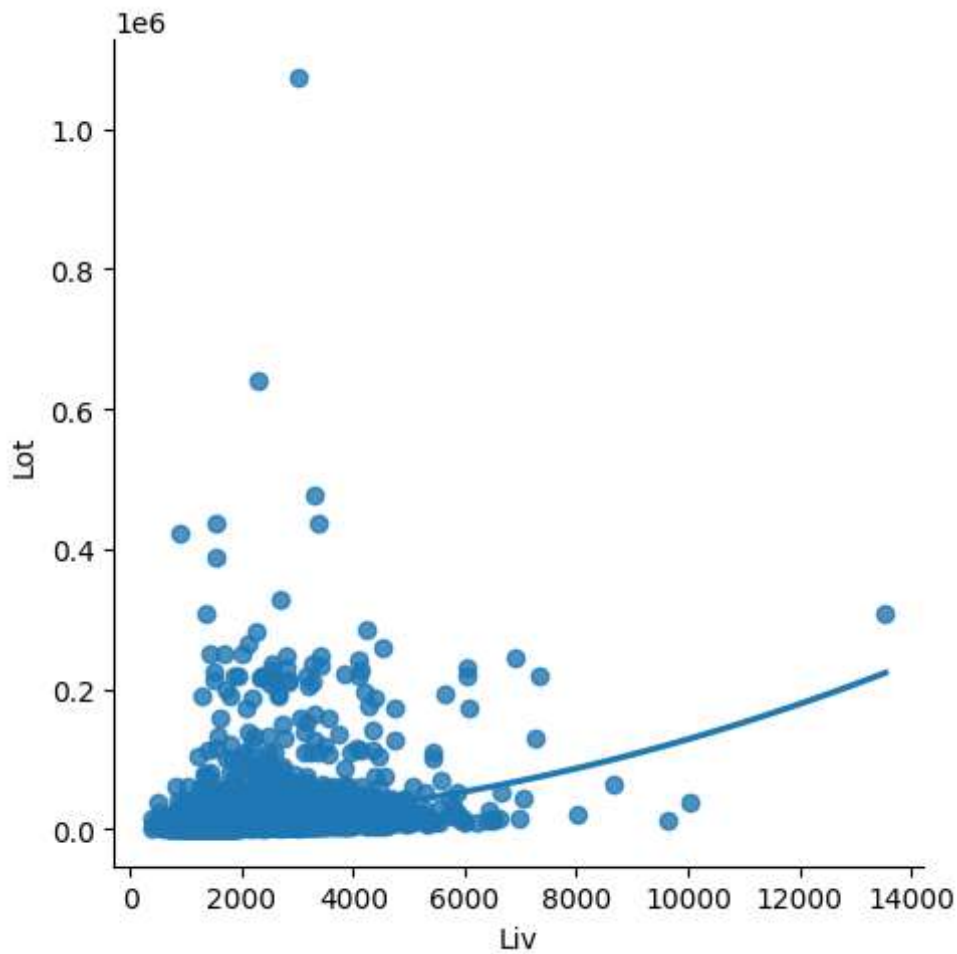
	Liv	Lot
0	1340	7912
1	3650	9050
2	1930	11947
3	2000	8030
4	1940	10500
5	880	6380
6	1350	2560
7	2710	35868
8	2430	88426
9	1520	6200

In [5]:

```
1 sns.lmplot(x='Liv',y='Lot',data=dt,order=2,ci=None)
```

Out[5]:

&lt;seaborn.axisgrid.FacetGrid at 0x160d1b52b50&gt;



In [6]:

```
1 dt.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4600 entries, 0 to 4599
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0    Liv    4600 non-null   int64  
 1    Lot    4600 non-null   int64  
dtypes: int64(2)
memory usage: 72.0 KB
```

In [7]:

```
1 dt.describe()
```

Out[7]:

	Liv	Lot
count	4600.000000	4.600000e+03
mean	2139.346957	1.485252e+04
std	963.206916	3.588444e+04
min	370.000000	6.380000e+02
25%	1460.000000	5.000750e+03
50%	1980.000000	7.683000e+03
75%	2620.000000	1.100125e+04
max	13540.000000	1.074218e+06

In [8]:

```
1 dt.fillna(method='ffill')
```

Out[8]:

	Liv	Lot
0	1340	7912
1	3650	9050
2	1930	11947
3	2000	8030
4	1940	10500
...	...	...
4595	1510	6360
4596	1460	7573
4597	3010	7014
4598	2090	6630
4599	1490	8102

4600 rows × 2 columns

In [9]:

```
1 x=np.array(dt['Liv']).reshape(-1,1)
2 y=np.array(dt['Lot']).reshape(-1,1)
```

In [10]:

```
1 dt.dropna(inplace=True)
```

C:\Users\HP\AppData\Local\Temp\ipykernel\_10920\735218168.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
dt.dropna(inplace=True)
```


In [11]:

```
1 X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
2 reg=LinearRegression()
3 reg.fit(X_train,y_train)
4 print(reg.score(X_test,y_test))
```

0.052631207637805266

In [12]:

```
1 y_pred=reg.predict(X_test)
2 plt.scatter(X_test,y_test,color='b')
3 plt.plot(X_test,y_pred,color='k')
4 plt.show()
```

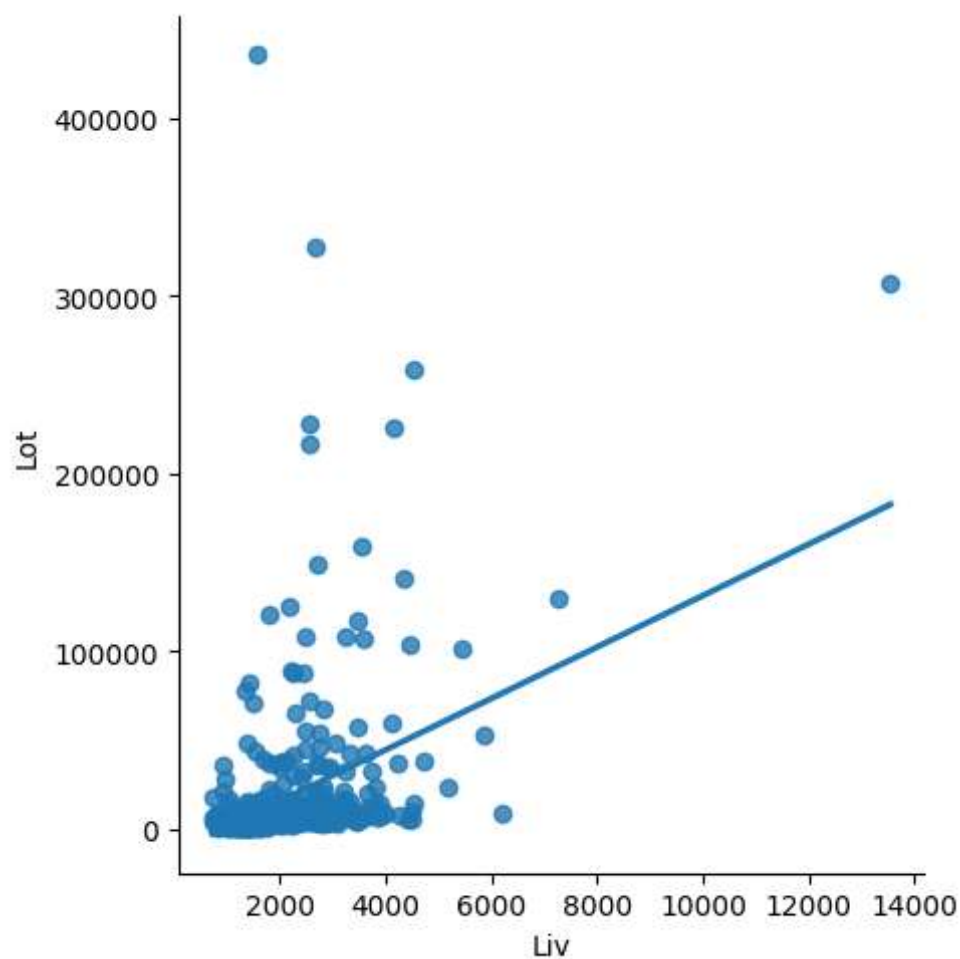


In [13]:

```
1 dt500=dt[:][:500]
2 sns.lmplot(x="Liv",y="Lot",data=dt500,order=1,ci=None)
```

Out[13]:

<seaborn.axisgrid.FacetGrid at 0x160c688ce90>




In [14]:

```

1 dt500.fillna(method='ffill',inplace=True)
2 X=np.array(dt500['Liv']).reshape(-1,1)
3 y=np.array(dt500['Lot']).reshape(-1,1)
4 dt500.dropna(inplace=True)
5 X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
6 reg=LinearRegression()
7 reg.fit(X_train,y_train)
8 print("Regression:",reg.score(X_test,y_test))
9 y_pred=reg.predict(X_test)
10 plt.scatter(X_test,y_test,color="b")
11 plt.plot(X_test,y_pred,color='k')
12 plt.show()

```

Regression: 0.11324919943741196



In [15]:

```

1 from sklearn.linear_model import LinearRegression
2 from sklearn.metrics import r2_score
3 model=LinearRegression()
4 model.fit(X_train,y_train)
5 y_pred=model.predict(X_test)
6 r2=r2_score(y_test,y_pred)
7 print("R2 score:",r2)

```

R2 score: 0.11324919943741196

In [16]:


```
1 from sklearn.linear_model import Ridge, RidgeCV, Lasso
2 from sklearn.preprocessing import StandardScaler
```

In [17]:

```
1 plt.figure(figsize = (10, 10))
2 sns.heatmap(dt.corr(), annot = True)
3
```

Out[17]:

&lt;Axes: &gt;





In [18]:

```
1 features = dt.columns[0:2]
2 target = dt.columns[-1]
3 #X and y values
4 X = dt[features].values
5 y = dt[target].values
6 #split
7 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
8 print("The dimension of X_train is {}".format(X_train.shape))
9 print("The dimension of X_test is {}".format(X_test.shape))
10 #Scale features
11 scaler = StandardScaler()
12 X_train = scaler.fit_transform(X_train)
13 X_test = scaler.transform(X_test)
14
```

The dimension of X\_train is (3220, 2)

The dimension of X\_test is (1380, 2)

In [19]:

```
1 #Model
2 lr = LinearRegression()
3 #Fit model
4 lr.fit(X_train, y_train)
5 #predict
6 #prediction = lr.predict(X_test)
7 #actual
8 actual = y_test
9 train_score_lr = lr.score(X_train, y_train)
10 test_score_lr = lr.score(X_test, y_test)
11 print("\nLinear Regression Model:\n")
12 print("The train score for lr model is {}".format(train_score_lr))
13 print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 1.0

The test score for lr model is 1.0

In [20]:

```
1 #Model
2 lr = LinearRegression()
3 #Fit model
4 lr.fit(X_train, y_train)
5 #predict
6 #prediction = lr.predict(X_test)
7 #actual
8 actual = y_test
9 train_score_lr = lr.score(X_train, y_train)
10 test_score_lr = lr.score(X_test, y_test)
11 print("\nLinear Regression Model:\n")
12 print("The train score for lr model is {}".format(train_score_lr))
13 print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 1.0  
The test score for lr model is 1.0

In [21]:

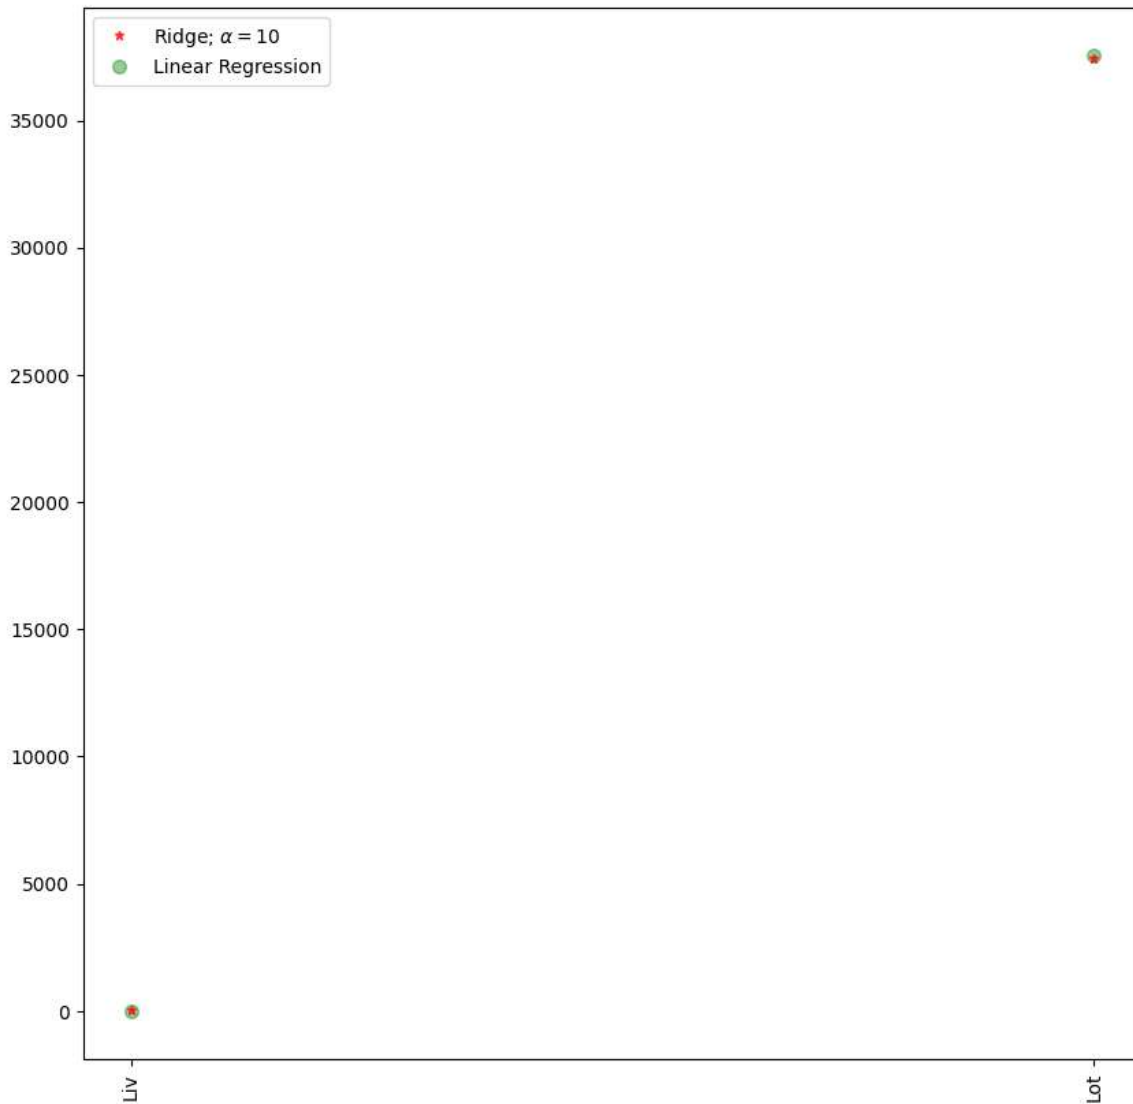
```
1 #Ridge Regression Model
2 ridgeReg = Ridge(alpha=10)
3 ridgeReg.fit(X_train,y_train)
4 #train and test scorefor ridge regression
5 train_score_ridge = ridgeReg.score(X_train, y_train)
6 test_score_ridge = ridgeReg.score(X_test, y_test)
7 print("\nRidge Model:\n")
8 print("The train score for ridge model is {}".format(train_score_ridge))
9 print("The test score for ridge model is {}".format(test_score_ridge))
10
```

Ridge Model:

The train score for ridge model is 0.9999900245012017  
The test score for ridge model is 0.9999902306741419

In [22]:

```
1 plt.figure(figsize = (10, 10))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red')
3 #plt.plot(rr100.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue')
4 plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
5 plt.xticks(rotation = 90)
6 plt.legend()
7 plt.show()
```



In [23]:

```
1 #Lasso regression model
2 print("\nLasso Model: \n")
3 lasso = Lasso(alpha = 10)
4 lasso.fit(X_train,y_train)
5 train_score_ls =lasso.score(X_train,y_train)
6 test_score_ls =lasso.score(X_test,y_test)
7 print("The train score for ls model is {}".format(train_score_ls))
8 print("The test score for ls model is {}".format(test_score_ls))
9
```

Lasso Model:

The train score for ls model is 0.9999999291360324


The test score for ls model is 0.9999999291231869

In [24]:

```
1 pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[24]:

&lt;Axes: &gt;



In [25]:

```
1 #Using the linear CV model
2 from sklearn.linear_model import LassoCV
3 #Lasso Cross validation
4 lasso_cv = LassoCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10], random_state=0).fit(X_train, y_train)
5 #score
6 print(lasso_cv.score(X_train, y_train))
7 print(lasso_cv.score(X_test, y_test))
8
```

0.9999999999997705


0.9999999999996928

In [26]:

```

1 #plot size
2 plt.figure(figsize = (10, 10))
3 #add plot for ridge regression
4 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red')
5 #add plot for lasso regression
6 #plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue')
7 #add plot for linear model
8 plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
9 #rotate axis
10 plt.xticks(rotation = 90)
11 plt.legend()
12 plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
13 plt.show()
14

```



In [27]:

```
1 #Using the linear CV model
2 from sklearn.linear_model import RidgeCV
3 #Ridge Cross validation
4 ridge_cv = RidgeCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10]).fit(X_train, y_train)
5 #score
6 print("The train score for ridge model is {}".format(ridge_cv.score(X_train, y_train))
7 print("The train score for ridge model is {}".format(ridge_cv.score(X_test, y_test)))
```

The train score for ridge model is 0.9999999999999996

The train score for ridge model is 0.9999999999999994

In [28]:

```
1 from sklearn.linear_model import ElasticNet
2 regr=ElasticNet()
3 regr.fit(X,y)
4 print(regr.coef_)
5 print(regr.intercept_)
```

[7.84570096e-07 9.99999995e-01]

-0.0016010777617339045

In [29]:

```
1 y_pred_elastic=regr.predict(X_train)
2
```

In [30]:

```
1 mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
2 print("Mean Squared Error on test set",mean_squared_error)
```

Mean Squared Error on test set 1635484254.2408667

In [ ]:

```
1
```